

Comparison of data processing on cloud using Spark with S3 and Redshift

Dublin City University
School of Computing
CA674: Cloud Architecture
Project report
Date of submission – 16th Dec 2018
Course – MCM

Name	Student Id no	Email id
Aishwarya Gupta	18210298	aishwarya.gupta3@mail.dcu.ie
Apurva Gawad	18210295	apurva.gawad2@mail.dcu.ie
Gautam Shanbhag	18210455	gautam.shanbhag2@mail.dcu.ie

Contents

Introduction.....	4
Amazon Redshift	4
Redshift Architecture:	4
Amazon Redshift Setup:.....	5
Spark	7
Spark Framework and Architecture:.....	7
Spark setup on Databricks:	8
Amazon S3 - Simple Storage Service.....	9
Amazon S3 Setup:.....	9
Comparative Analysis of data processing capabilities of Amazon Redshift and Spark with S3 ...	10
Implementation Plan:	10
RedShift Readings:	10
Data Loading from S3 to Redshift	10
DDL Commands:	11
DML commands:.....	13
Unloading data from Redshift to S3:	15
Spark Readings:	16
Data Loading from S3 to Spark:	16
DDL Commands:	16
DML Commands:	17
Data unloading to S3 from Spark:.....	18
Summary of the review above for some parameters	19
Conceptual Comparison:-	20
Conclusion:	20
References:	21

Figure 1. Redshift Architecture	5
Figure 2. Step 2. for Redshift Setup	6
Figure 3. Step 6. Cluster Dashboard.....	6
Figure 4. Spark Framework	7
Figure 5. Spark architecture	8
Figure 6. Cluster configuration	9
Figure 7. Step 8. Setting permissions for S3	10
Figure 8. Redshift data loading from S3.....	11
Figure 9. Redshift Alter command	12
Figure 10. Redshift Drop Command.....	13
Figure 11. Redshift Update command	14
Figure 12. Redshift Nested query command	15
Figure 13. Redshift Unload command	15
Figure 14. Spark data loading from S3	16
Figure 15. Spark Delete command.....	16
Figure 16. Spark Aggregate Query	17
Figure 17. Spark unload command	18

Introduction

Cloud computing is the on-demand delivery of compute power, database storage, applications, and other IT resources through a cloud services platform via the internet with pay-as-you-go pricing [5]. Cloud has opened gateways to big data processing due to its distributed framework. In this project, we will be comparing PaaS systems Spark (using data storage as S3) on Databricks, which is an open source cluster-computing framework for data processing with Redshift, which is a data ware housing system, which stores and analyses large amounts of data.

Amazon Redshift

Amazon Redshift is a data warehouse cloud service that helps companies to store and analyse large amounts of data, up to the petabyte scale. It forms a part of the large cloud-computing platform Amazon Web Services. It is built on top of technology from the massive parallel processing (MPP) data warehouse company ParAccel (later acquired by Actian), to handle large-scale data sets and database migrations. It is an enterprise-class relational database query and management system. It is based on a PostgreSQL 8.0.2, but has made many modifications in that version. It makes use of JDBC and ODBC connections to interact with other applications. It supports client connections with many types of applications, including business intelligence (BI), reporting, data, and analytics tools. There is a consensus that high performance and low costs are key advantages of using Amazon Redshift. Amazon Redshift achieves efficient storage and optimum query performance through a combination of massively parallel processing, columnar data storage, and very efficient, targeted data compression encoding schemes

Redshift Architecture:

Amazon Redshift works on the premise of collecting several nodes and creating an Amazon Redshift cluster. Clients can provision this cluster and upload data, then use complex data analysis queries to obtain business intelligence analysis

The client applications interact with the data warehouse cluster using JDBC or ODBC connections through the leader node. The leader node develops execution plans to carry out database operations. Based on the execution plan, the leader node compiles code, distributes the compiled code to the compute nodes, and assigns a portion of the data to each compute node. The leader node distributes SQL statements to the compute nodes only when a query references tables that are stored on the compute nodes. All other queries run exclusively on the leader node. The compute nodes execute the compiled code and send intermediate results back to the leader node for final aggregation. A compute node is partitioned into node slices. Each slice is allocated a portion of the node's memory and disk space, where it processes a portion of the workload assigned to the node. A cluster contains one or more databases.

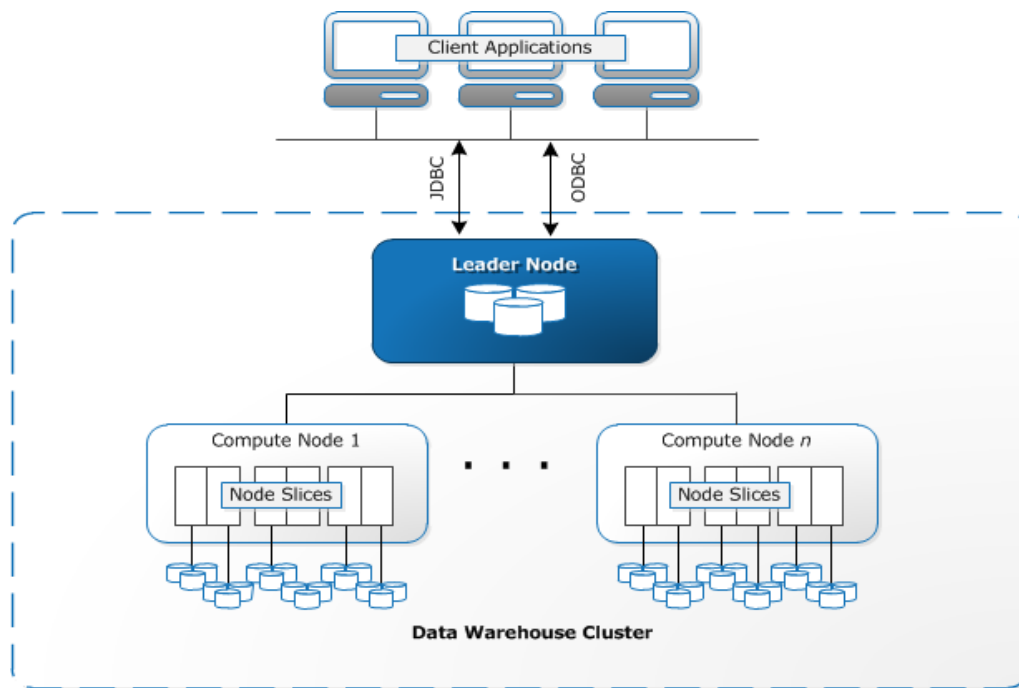


Figure 1. Redshift Architecture

Amazon Redshift is a relational database management system (RDBMS), so it is compatible with other RDBMS applications. Although it provides the same functionality as a typical RDBMS, including online transaction processing (OLTP) functions such as inserting and deleting data, Amazon Redshift is used for Analytical processing (OLAP), it is optimized for high-performance analysis and reporting of very large datasets.

Amazon Redshift Setup:

1. Before starting with Amazon Redshift Configuration, we need to setup an AWS account.
2. After Setting up AWS, you need to complete IAM configurations (Identity and Access Management) AWS Dashboard → Security, Identity & Compliance → IAM
3. Create Roles for Redshift and S3 Access Using IAM role setting. CLICK Roles → Create Role → Select Redshift → Select Case as Redshift - Customizable → Next: Permissions → Check 1. AmazonRedshiftFullAccess & 2. AmazonS3FullAccess → Next: Permissions → Tags (Optional) Next: Review
4. Newly Created roles will appear on the IAM Console
5. Create Cluster using Redshift → Click on Launch Cluster
6. Provide cluster name, provide DB name and password (by default DB name: dev)
7. Select Node as dc2.large, Cluster type: Single (Free others are paid) → Continue
8. Basic Configuration → Select Available IAM role to the one previously created with permissions → Click Continue

Redshift dashboard

Launch your Amazon Redshift cluster - Advanced settings | [Switch to quick launch](#)

CLUSTER DETAILS NODE CONFIGURATION ADDITIONAL CONFIGURATION REVIEW

Provide the details of your cluster. Fields marked with * are required.

Cluster identifier* This is the unique key that identifies a cluster. This parameter is stored as a lowercase string. (e.g. my-dw-instance)

Database name Optional. A default database named dev is created for the cluster. Optionally, specify a custom database name (e.g. mydb) to create an additional database.

Database port* Port number on which the database accepts connections.

Master user name* Name of master user for your cluster. (e.g. awsuser)

Master user password* Password must contain 8 to 64 printable ASCII characters excluding /, ", \, and @. It must contain 1 uppercase letter, 1 lowercase letter, and 1 number.

Confirm password* Confirm master user password

Figure 2. Step 2. for Redshift Setup

9. Review Configurations → Click Launch Cluster
10. The cluster would be created successfully and would appear on dashboard with status available and healthy.
11. We will connect with this Redshift cluster from our local system using DBVisualizer tool.

Redshift dashboard

Clusters

Query editor New

Saved queries

Snapshots

Security

Parameter groups

Workload management

Reserved nodes

Advisor

Events

Connect client

What's new

Announcing availability of Elastic Resize.
Using Elastic Resize, you can now add or subtract nodes to your Amazon Redshift cluster within minutes. You can do elastic resize using the Amazon Redshift console or by running the `resize-cluster` CLI command. Your cluster version must be 1.0.4652 or higher. [Learn more](#)

	Cluster	Cluster Status	DB Health	Release Status	In Maintenance	Recent Events	Config timeline
<input type="checkbox"/>	training1	available	healthy	Up to date	no	6	View timeline

Figure 3. Step 6. Cluster Dashboard

Spark

Spark is an open source unified analytics engine maintained by Apache software Foundation for large-scale data processing. It is a general-purpose cluster computing framework. On top of the Spark core data processing engine, there are libraries for SQL, machine learning, graph computation, and stream processing, which can be used together in an application. Application developers and data scientists incorporate Spark into their applications to rapidly query, analyse, and transform data at scale. Tasks most frequently associated with Spark include ETL and SQL batch jobs across large data sets, processing of streaming data from sensors, IoT, or financial systems, and machine learning tasks.

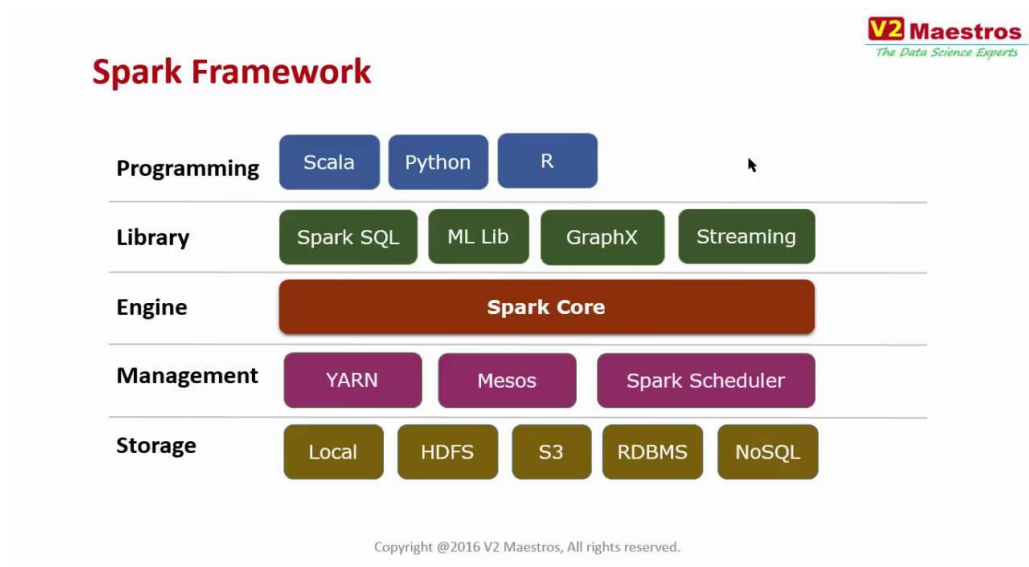


Figure 4. Spark Framework

Spark Framework and Architecture:

Spark uses a Master/Worker architecture, the driver works as the master and the executors work as the slaves.

The driver is a JVM process that holds the SparkContext for a spark application, it splits the incoming tasks and schedules them to run on executors. It coordinates between the workers and manages the overall schedule of tasks.

Executor is a distributed agent that is responsible for the execution of a task. When an executor begins, it communicates directly with the driver to execute tasks.

RDDs or Resilient distributed datasets are the back-bone spark core, it is the primary data structure that is used by spark. A RDD is a resilient and distributed collection of records spread over one or many partitions. Data frames that are used in SparkSQL are internally processed as RDDs but provide a schema to the data set.

Spark SQL is apache library for structured data processing. Unlike the basic Spark RDD API, the interfaces provided by Spark SQL provide Spark with more information about the structure of both the data and the computation being performed [4]. One use of Spark SQL is to execute SQL queries.

Spark can be used in two ways, either by running it on your local machine, where the JVM holds both the driver and executor programs, or in cluster mode – where the executors run in the cluster and you have the freedom to choose where to run the driver. Either on your JVM (generally in case of interactive notebooks) or on the cluster (in case of batch processing or streaming applications).

Spark needs a cluster manager to work as you can see in the figure 1, the cluster manager can either be an external manager like Mesos or Hadoop YARN, or we can use spark's standalone manager that comes with spark. Selection of a cluster manager for a spark application is dependent on the purpose of the application because every cluster manager provides a different set of scheduling capabilities.




Figure 5. Spark architecture

Unlike other data processing engines e.g. Apache Hadoop or Redshift, Spark does not have an internal storage system which in a way gives the user the freedom to use multiple storage systems like S3, HDFS, etc. This makes spark more flexible, however also gives additional burden of storage maintenance.

For this project, we will be using Apache Spark with S3 data storage, and executing it on cluster using Databricks community Edition.

Spark setup on Databricks:

Spark is a framework that can be used over a cluster using interactive notebooks. For our project we will be using Spark over a cluster on Databricks notebook. Databricks is a company founded by the creators of Apache Spark, that aims to help clients with cloud-based big data processing using Spark [6]. It provides a platform on which we can run spark for various analytical purposes.

1. Create a Databricks account and login to view the Dashboard
2. Click on Cluster → Create Cluster
3. Provide Cluster Name, Spark Version, Python Version & Zone → Create Cluster

Figure 6. Cluster configuration

4. Created Cluster would be visible on Cluster Page and Status will change from pending to Running
5. On Dashboard → Click Create Notebook, provide Name, Programming language (we are using Python) and select cluster from the dropdown.
6. The notebook can be used for unified analytics processing using spark.

Amazon S3 - Simple Storage Service

Amazon Simple Storage Service (Amazon S3) is an object storage service that offers industry-leading scalability, data availability, security, and performance. This means customers of all sizes and industries can use it to store and protect any amount of data for a range of use cases, such as websites, mobile applications, backup and restore, archive, enterprise applications, IoT devices, and big data analytics. Amazon S3 provides easy-to-use management features so you can organize your data and configure finely-tuned access controls to meet your specific business, organizational, and compliance requirements [3].

Amazon S3 Setup:

1. Before starting with Amazon S3 Configuration, we need to setup an AWS account.
2. After Setting up AWS, you need to complete IAM configurations (Identity and Access Management) AWS Dashboard → Security, Identity & Compliance → IAM
3. Create Roles for Redshift and S3 Access Using IAM role setting. CLICK Roles → Create Role → Select Redshift → Select Case as Redshift - Customizable → Next: Permissions → Check 1. AmazonRedshiftFullAccess & 2. AmazonS3FullAccess → Next: Permissions → Tags (Optional) Next: Review

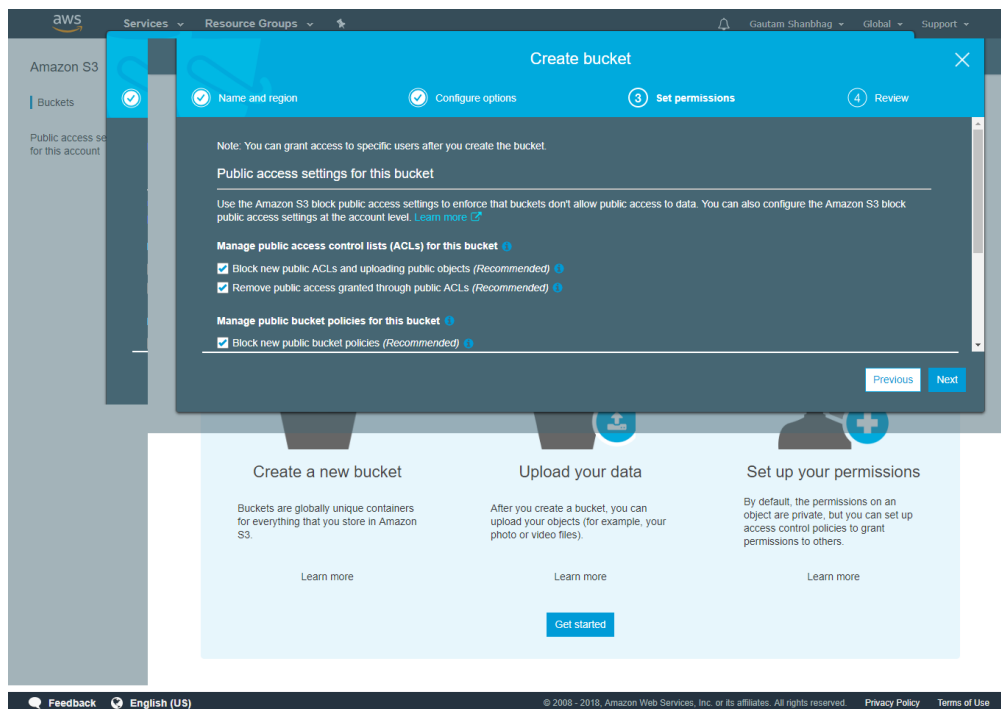


Figure 7. Step 8. Setting permissions for S3

4. Newly Created roles will appear on the IAM Console.
5. Select S3 from Dashboard → Click on Create Bucket.
6. Provide a unique relevant bucket name and select the closest geographical region.
7. Basic configuration → Click Next.
8. Set Permissions → Click Next.
9. Click on Create bucket.
10. S3 Created bucket would appear on the S3 Dashboard

Comparative Analysis of data processing capabilities of Amazon Redshift and Spark with S3

Implementation Plan:

After completing the setup, we are ready to load our data and perform ETL and Analytical processing of data. For our analysis we will be using a large data set which we have acquired from Kaggle.com. This data set contains information about transactions made in a retail store on Black Friday Sale. The data set contains 550,000 observations.

First, we need to load our data on S3 bucket, we will load data into Redshift and Spark from our S3 bucket.

RedShift Readings:

Data Loading from S3 to Redshift:

1. To load the data from S3, we need to first create a table structure. Since, we need to create a table structure it implies that we should be previously aware of the type of data we are loading and cannot be used for Dynamic loading of data.

2. After creating the table, we will use the COPY command to load the data from S3 to Redshift.
3. Query used:

```
copy Black_Friday from
's3://trainingashbucket/BlackFriday.csv' iam_role
'arn:aws:iam::319165558774:role/TrainingAsh' delimiter
',' IGNOREHEADER 1;
```

4. Time consumed: 50.238 s

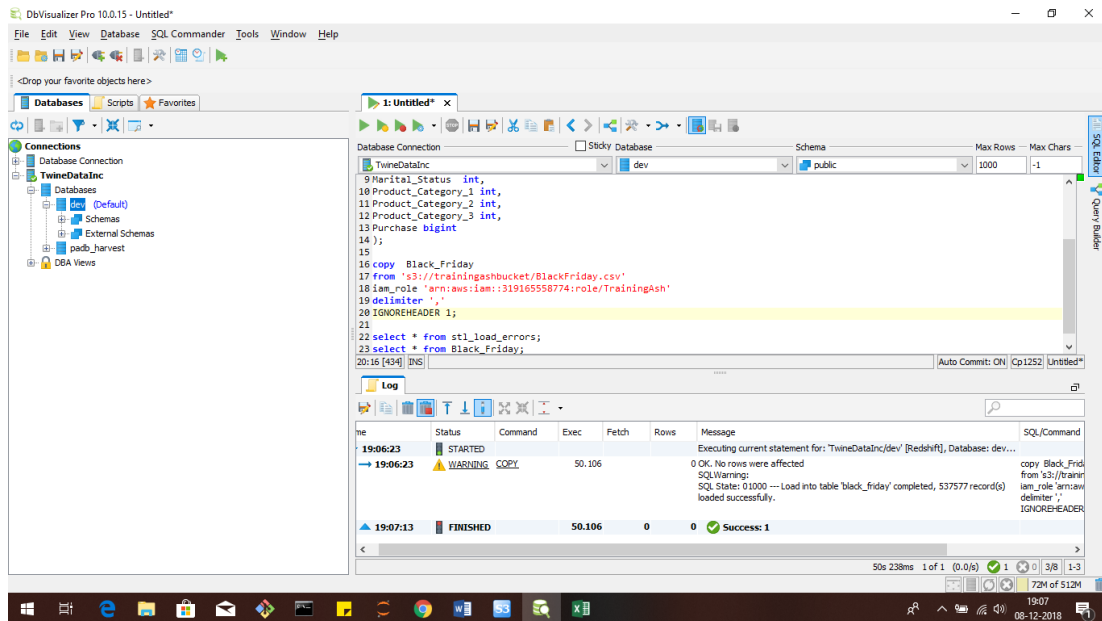


Figure 8. Redshift data loading from S3

DDL Commands:

- 1) Create table command:

- a) Query used:

```
create table Black_Friday(
User_ID bigint, Product_ID varchar, Gender char(1),
Age varchar, Occupation int, City_Category
char(1), Stay_In_Current_City_Years varchar,
Marital_Status int, Product_Category_1 int,
Product_Category_2 int, Product_Category_3 int,
Purchase bigint );
```

- b) First Execution time: 474 ms

- c) Average Time: 172 ms

- 2) Alter table command: We created another column purchase_in_euro which stored the purchase amount in Euros.

- a) Query used:

```
alter table Black_Friday
add column purchase_in_euro bigint;
```

b) Execution time: 3s 607ms

c) Average Time: 172 ms

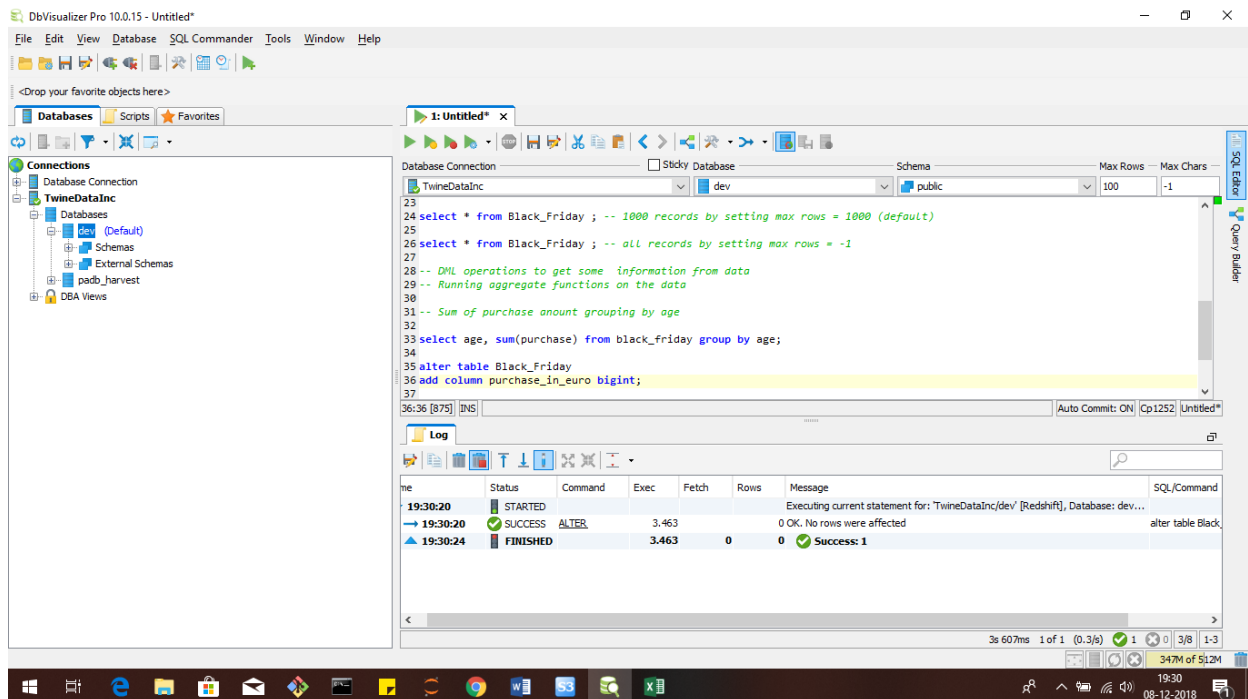


Figure 9. Redshift Alter command

- 3) Delete command: We use delete command with various conditions to test how much of its performance is affected by using multiple conditions. Since, Redshift is a columnar database, such conditions increase the computation time.

a) Query used:

```
delete from black_friday where
(black_friday.product_category_1 is null or
black_friday.product_category_2 is null or
black_friday.product_category_3 is null) and
black_friday.city_category != 'A' ;
```

b) Exec time: 2s 183ms

c) Average time: 357 ms

- 4) Drop command: Drop command in Redshift frees up the memory when executed. This command is very fast owing to its multiple parallel processing and columnar data storage structure.

a) Query used:

```
drop table Black_Friday;
```

b) Exec time: 162 ms

c) Average time: 153 ms

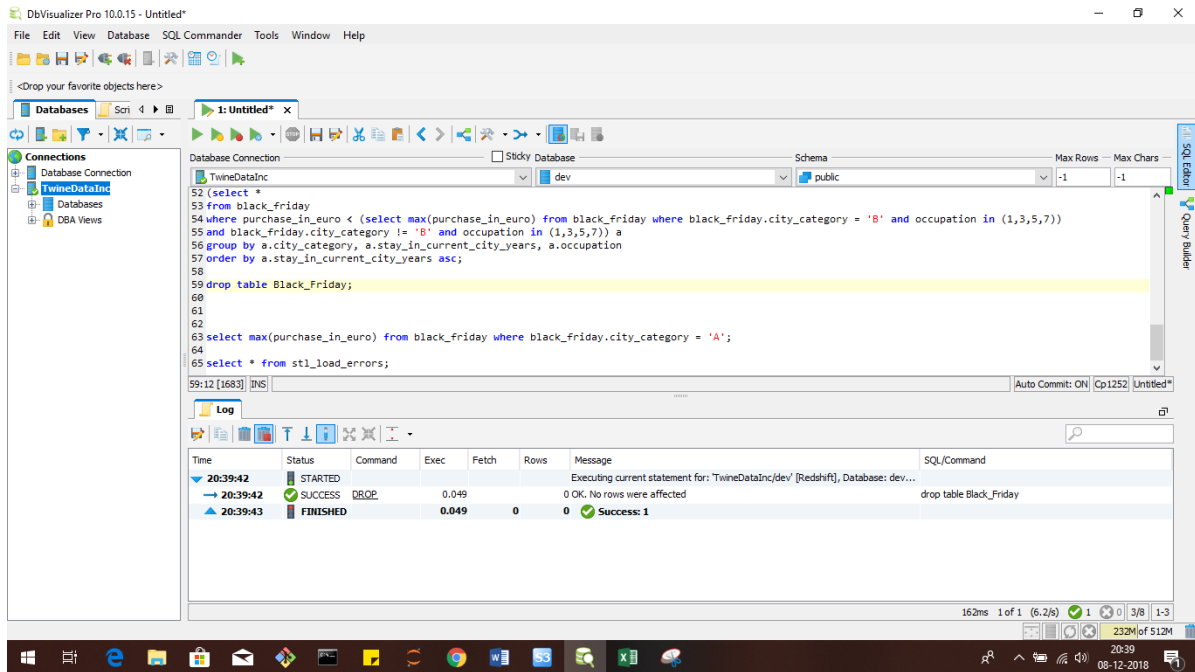


Figure 10. Redshift Drop Command

DML commands:

- 1) Selecting 1000 records from table : We especially ran two queries to see if there is any major difference in the time consumed to fetch fraction of data and complete data from Redshift table.

a. Query used:

```
select * from Black_Friday limit 1000;
```

b. Execution time: 3m 708s

c. Average execution time: 480 ms

- 2) Fetching entire data set of 537577 records.

a. Query Used:

```
select * from Black_Friday
```

b. Execution time: 21 s 257ms

c. Average execution time: 3s 454 ms

- 3) Update Query: We wanted to insert data in the column that we had created using the alter command. Since this is a column level function, we expected this command to perform well.

a. Query used:

```
update Black_Friday set purchase_in_euro = purchase *
0.88;
```

b. Exec time: 2s 691ms

c. Average exec time : 1s 290ms

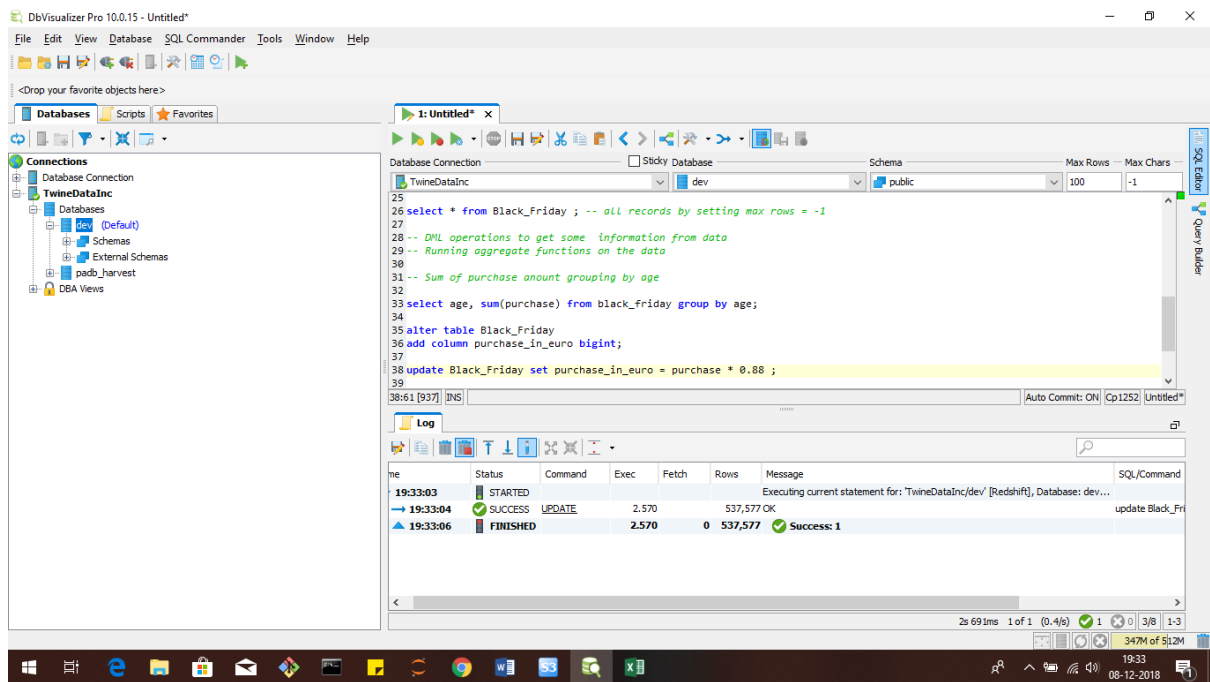


Figure 11. Redshift Update command

- 4) Aggregate functions: Redshift being a columnar database is fast in terms of aggregate functions, hence this was an important comparative factor for us with spark.

a. Query used:

```

select    gender,    age,    avg(purchase_in_euro)    as
average_spent    from    black_friday    where
black_friday.marital_status = 1 group by gender, age
having avg(purchase_in_euro) > 1000 order by age asc,
average_spent asc;

```

b. Exec time: 8s 193 ms

c. Average time: 393 ms

- 5) Nested Query: We wanted to see how Redshift performs when nested queries are applied, we have used three levels nest to see the performance of Redshift.

a. Query used:

```

select  a.city_category,  a.stay_in_current_city_years,
a.occupation,  max(a.purchase_in_euro),  a.gender from
(select * from black_friday where purchase_in_euro <
(select max(purchase_in_euro) from black_friday where
black_friday.city_category = 'C' and occupation in
(2,4,5,7)) and black_friday.city_category != 'C' and
occupation in (2,4,5,7)) a group by a.city_category,
a.stay_in_current_city_years,  a.occupation,  a.gender
having    count(a.user_id)    >    2    order    by
a.stay_in_current_city_years asc;

```

b. Exec time: 30s 629 ms

c. Average execution time: 406 ms

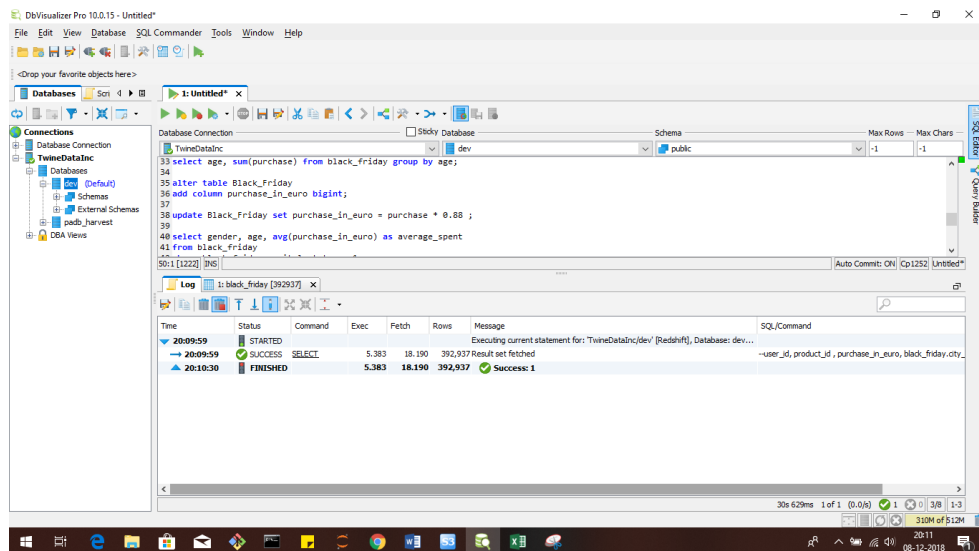


Figure 12. Redshift Nested query command

Unloading data from Redshift to S3:

One of the benefits of using Redshifts is that it is a data ware house hence, we can store are processed data in Redshift during the Loading of the ETL process. However, when using Spark, since it does not have its own storage we need to unload the data from spark and export or load it to another data storage system which is S3 for our case. To compare the performance of data unloading between spark and S3 we have unloaded data from Redshift to S3.

1) Query Used:

```
unload('select * from Black_friday') to
's3://trainingashbucket/Black_friday_output_redshift/'
iam_role 'arn:aws:iam::319165558774:role/TrainingAsh'
delimiter as ',' parallel off;
```

2) Exec time : 7s 298 ms

3) Average time : 1s 578ms

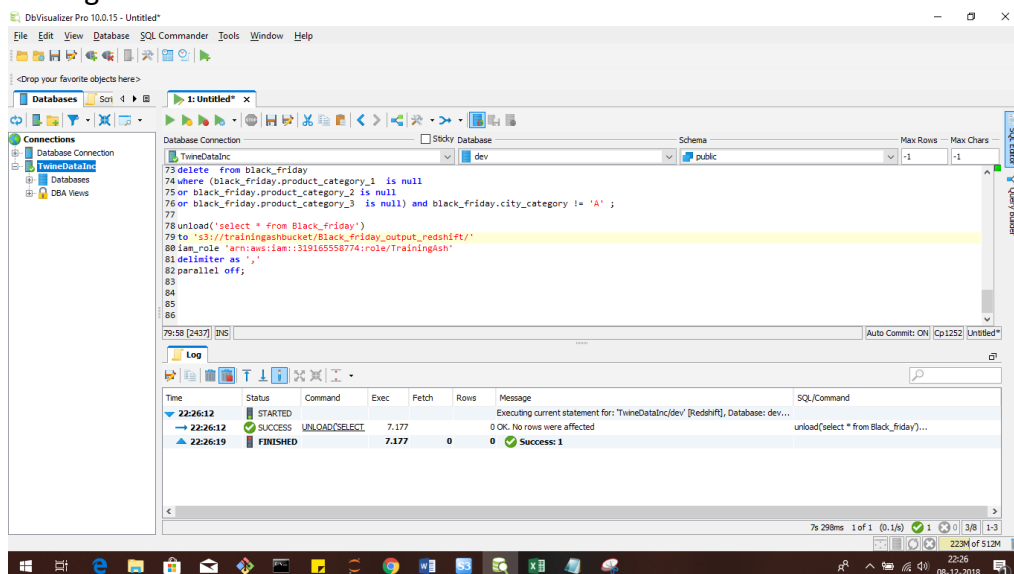


Figure 13. Redshift Unload command

Spark Readings:

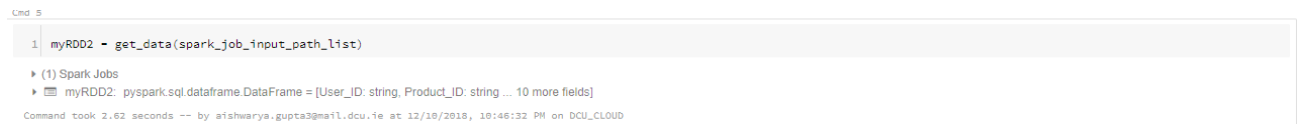
Data Loading from S3 to Spark:

- 1) Data can be loaded into spark without creating any previous table. The dataframe gets created dynamically at run time when the data is being loaded.

- 2) Exec Time: 2.62s

- 3) Command used:

```
spark.read.csv(spark_job_input_path, header = "TRUE")
```



```
Cmd 5
1 myRDD2 = get_data(spark_job_input_path_list)

(1) Spark Jobs
myRDD2: pyspark.sql.dataframe.DataFrame = [User_ID: string, Product_ID: string ... 10 more fields]
Command took 2.62 seconds -- by aishwarya.gupta3@mail.dcu.ie at 12/10/2018, 10:46:32 PM on DCU_CLOUD
```

Figure 14. Spark data loading from S3

DDL Commands:

- 1) Alter Table command – As we are working with data frames, there is no direct DDL command to alter the data frame, however we can use spark commands to alter the data frames.

- a. Exec time – 0.02 s

- b. Command used:


```
myRDD = myRDD.withColumn('purchase_in_euro',lit(0))
```

- 2) Delete command – Spark does not have a direct mechanism to delete rows based on condition, however we can select the rows that we want and create a new data frame replacing the old one.

- a. Exec time – 0.11s

- b. Command used:

```
myRDD =
myRDD.where((col('Product_Category_1').isNotNull()) &
(col('Product_Category_2').isNotNull()) &
(col('Product_Category_3').isNotNull()) |
(col('City_Category') == 'A'))
```



```
Cmd 19
1 myRDD = myRDD.where((col('Product_Category_1').isNotNull()) & (col('Product_Category_2').isNotNull()) & (col('Product_Category_3').isNotNull()) | (col('City_Category') == 'A'))

myRDD: pyspark.sql.dataframe.DataFrame = [User_ID: string, Product_ID: string ... 11 more fields]
Command took 0.11 seconds -- by aishwarya.gupta3@mail.dcu.ie at 12/8/2018, 10:39:14 PM on DCU_CLOUD
```

Figure 15. Spark Delete command

- 3) Drop command – In Spark there is no way to directly drop the table, however the reference could be changed and cache cleared to get the same result as that of drop.

DML Commands:

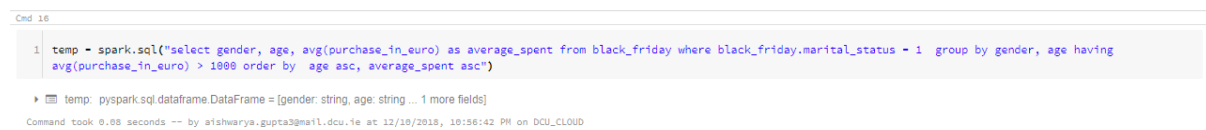
- 1) Fetching 1000 rows without any condition
 - a. Exec time – 0.03 s
 - b. Command used:

```
temp = spark.sql("select * from Black_Friday limit 1000")
```
- 2) Fetching 1000 rows without any condition – We can see that there is no major difference in the time consumed by spark to fetch the entire data and part of the data compared to Redshift.
 - a. Exec time – 0.03 s
 - b. Command used:

```
temp = spark.sql("select * from Black_Friday")
```
- 3) Update Query: Since we are working on data frame, we cannot update the columns directly using sql commands, however we can manipulate the data frame and load values.
 - a. Exec time – 0.03 s
 - b. Command used:

```
myRDD = myRDD.withColumn('purchase_in_euro', col('purchase')*0.88)
```
- 4) Aggregate Query: Even though, Redshift is a columnar data base, we can see here that spark is clearly faster when it comes to aggregate queries. This is due to the portioning of work mechanism that spark uses.
 - a. Exec time – 0.08 s
 - b. Command used:

```
temp = spark.sql("select gender, age, avg(purchase_in_euro) as average_spent from black_friday where black_friday.marital_status = 1 group by gender, age having avg(purchase_in_euro) > 1000 order by age asc, average_spent asc")
```



```
1 temp = spark.sql("select gender, age, avg(purchase_in_euro) as average_spent from black_friday where black_friday.marital_status = 1 group by gender, age having avg(purchase_in_euro) > 1000 order by age asc, average_spent asc")
```

temp: pyspark.sql.dataframe.DataFrame = [gender: string, age: string ... 1 more fields]

Command took 0.08 seconds -- by aishwarya.gupta3@gmail.com at 12/10/2018, 18:56:42 PM on DCU_CLOUD

Figure 16. Spark Aggregate Query

- 5) Nested Query : We can see that spark is extremely fast in nested queries.
 - a. Exec time – 0.10s
 - b. Command used:

```
temp = spark.sql("select a.city_category, a.stay_in_current_city_years, a.occupation, max(a.purchase_in_euro), a.gender from (select * from black_friday where purchase_in_euro < (select max(purchase_in_euro) from black_friday where black_friday.city_category = 'C' and occupation in (2,4,5,7)) and black_friday.city_category != 'C' and
```

```
occupation in (2,4,5,7)) a group by a.city_category,  
a.stay_in_current_city_years, a.occupation, a.gender  
having count(a.user_id) > 2 order by  
a.stay_in_current_city_years asc ")
```

Data unloading to S3 from Spark:

As Spark does not have its own data base or storage, we have to export our results to an external storage or else all work will be lost.

1) Exec time - 36.76 s

2) Command used:

```
myRDD.coalesce(1).write.save(OUTPUT_FILE,sep=",",format  
="csv",header="TRUE")
```



The screenshot shows a terminal window with two command prompts. The first prompt, labeled 'Cmd 21', shows the assignment of 'OUTPUT_FILE' to an S3 path. The second prompt, labeled 'Cmd 22', shows the Spark command to save the RDD to S3 in CSV format. Below the command, the execution time is shown as 36.76 seconds. The terminal also shows a Spark job summary indicating 1 job was executed.

```
Cmd 21  
1 OUTPUT_FILE = "s3n://trainingashbucket/Black_friday_output_spark_final"  
  
Command took 0.01 seconds -- by arishwarya.gupta3@mail.dcu.ie at 12/8/2018, 10:41:37 PM on DCU_CLOUD  
  
Cmd 22  
1 myRDD.coalesce(1).write.save(OUTPUT_FILE,sep=",",format="csv",header="TRUE")  
  
▶ (1) Spark Jobs  
Command took 36.76 seconds -- by arishwarya.gupta3@mail.dcu.ie at 12/8/2018, 10:41:42 PM on DCU_CLOUD  
fret 7%
```

Figure 17. Spark unload command

We can see that spark takes significant amount of time to export data; this is due to the fact that the RDDs are partitioned and collaborating the RDDs and sending the information takes time.

This phenomenon can also be observed when we try to show() the contents of the data frame, as again this requires the merging of RDDs which becomes time consuming for spark.

Summary of the review above for some parameters:-

The parameters for Redshift, which are in bracket, tells the first execution time of the query, it can be observed that repeated execution of same query tremendously reduces the execution time. This is due to the feature of result caching in Redshift. Despite result caching, we can observe from the readings that spark is significantly faster than Redshift when it comes to processing large amounts of data.

Comparison factor	Amazon Redshift	Apache Spark
Data Load time from S3	6s 663ms (50s 238 ms *)	2s 620 ms
Fetching entire data set	12 s 557ms (21 s 257ms *)	20 ms
Aggregate function execution	393 ms (8s 193 ms *)	80 ms
Nested query execution	406 ms (30s 629 ms *)	100 ms
Deletion of values based on condition	357 ms (2s 183 ms *)	110 ms
Drop table	153 ms (162 ms *)	10 ms
Alter command	185 ms (3s 607ms *)	20 ms
Update command	1s 290 ms (2s 691ms *)	30 ms
Data unload to S3	1s 578ms (7s 298 ms *)	36s 760 ms

Conceptual Comparison:-

In the below table we can see the comparison of various conceptual factors between Amazon Redshift and Apache spark.

Comparison factor	Amazon Redshift	Apache Spark
Category	Analytical database	Data processing engine
Supported languages	SQL-Like language	Extensible via pre built libraries in Java, Python, R, Scala
Interoperability	Compatible with standard JDBC, ODBC drivers Integrated with most BI, ETL tools	Runs Everywhere. Can access diverse data sources. Runs on Hadoop, Mesos, Kubernetes, standalone or in the cloud.
Distribution model	Commercial/managed	Opensource
Storage Support	Inbuilt Storage	External Storage Engine Required
Live app Database	Not supported	Supports streaming data
Concurrency	Efficient multi user support	Performance degrades as concurrency increases
Typical application	Data warehouse	applications

Conclusion:

Even though spark and Redshift are comparable when it comes to data processing, it all boils down to the requirement of the user. Redshift being a data ware house is preferable when the goal is to archive the data and requires periodical processing of large amount of data, having easy scalability it is suited for applications where the data is ever growing and of different types. However, spark comes handy when we need to process large amounts of data continuously; it is capable of handling streaming data and building machine learning algorithms.

References:

- [1] Online reference - <https://spark.apache.org/>
- [2] Image reference - <https://www.youtube.com/watch?v=ZTFGwQaXJm8>
- [3] <https://aws.amazon.com/s3/>
- [4] <https://spark.apache.org/docs/latest/sql-programming-guide.html>
- [5] <https://aws.amazon.com/what-is-cloud-computing/>
- [6] <https://databricks.com/product/unified-analytics-platform>
- [7] <https://dbengines.com/en/system/Amazon+Redshift%3BSpark+SQL>
- [8] <https://docs.aws.amazon.com/redshift/latest/mgmt/welcome.html>
- [9] <https://docs.aws.amazon.com/redshift/latest/mgmt/managing-clusters-console.html>
- [10] <https://docs.aws.amazon.com/redshift/latest/mgmt/generating-user-credentials.html>